

# GreenDroid: A Mobile Application Processor for a Future of Dark Silicon

---

Nathan Goulding, Jack Sampson, Ganesh Venkatesh,  
Saturnino Garcia, Joe Auricchio, Jonathan Babb<sup>+</sup>,  
Michael B. Taylor, and Steven Swanson

*Department of Computer Science and Engineering,  
University of California, San Diego*

<sup>+</sup> CSAIL, Massachusetts Institute of Technology

# **Where does dark silicon come from? (And how dark is it going to be?)**

Utilization Wall:

With each successive process generation, the percentage of a chip that can actively switch drops exponentially due to power constraints.

# We've Hit The Utilization Wall

*Utilization Wall: With each successive process generation, the percentage of a chip that can actively switch drops exponentially due to power constraints.*

- Scaling theory
  - Transistor and power budgets are no longer balanced
  - Exponentially increasing problem!
- Experimental results
  - Replicated a small datapath
  - More "dark silicon" than active
- Observations in the wild
  - Flat frequency curve
  - "Turbo Mode"
  - Increasing cache/processor ratio

# We've Hit The Utilization Wall

*Utilization Wall: With each successive process generation, the percentage of a chip that can actively switch drops exponentially due to power constraints.*

## ■ Scaling theory

- Transistor and power budgets are no longer balanced
- Exponentially increasing problem!

## ■ Experimental results

- Replicated a small datapath
- More "dark silicon" than active

## ■ Observations in the wild

- Flat frequency curve
- "Turbo Mode"
- Increasing cache/processor ratio

## Classical scaling

Device count	$S^2$
Device frequency	$S$
Device power (cap)	$1/S$
Device power ( $V_{dd}$ )	$1/S^2$
<b>Utilization</b>	<b>1</b>

## Leakage-limited scaling

Device count	$S^2$
Device frequency	$S$
Device power (cap)	$1/S$
<b>Device power (<math>V_{dd}</math>)</b>	<b><math>\sim 1</math></b>
<b>Utilization</b>	<b><math>1/S^2</math></b>

# We've Hit The Utilization Wall

*Utilization Wall: With each successive process generation, the percentage of a chip that can actively switch drops exponentially due to power constraints.*

## ■ Scaling theory

- Transistor and power budgets are no longer balanced
- Exponentially increasing problem!


## ■ Experimental results

- Replicated a small datapath
- More "dark silicon" than active

## ■ Observations in the wild

- Flat frequency curve
- "Turbo Mode"
- Increasing cache/processor ratio


Expected utilization for fixed area and power budget



# We've Hit The Utilization Wall

*Utilization Wall: With each successive process generation, the percentage of a chip that can actively switch drops exponentially due to power constraints.*

- **Scaling theory**
  - Transistor and power budgets are no longer balanced
  - Exponentially increasing problem!
- **Experimental results**
  - Replicated a small datapath
  - More "dark silicon" than active
- **Observations in the wild**
  - Flat frequency curve
  - "Turbo Mode"
  - Increasing cache/processor ratio



# We've Hit The Utilization Wall

*Utilization Wall: With each successive process generation, the percentage of a chip that can actively switch drops exponentially due to power constraints.*

## ■ Scaling theory


- Transistor and power budgets are no longer balanced
- Exponentially increasing problem!

## ■ Experimental results

- Replicated a small datapath
- More "dark silicon" than active

## ■ Observations in the wild

- Flat frequency curve
- "Turbo Mode"
- Increasing cache/processor ratio



# We've Hit The Utilization Wall

*Utilization Wall: With each successive process generation, the percentage of a chip that can actively switch drops exponentially due to power constraints.*

## ■ Scaling theory

Utilization @ 40 mm<sup>2</sup>, 3 W

The utilization wall will change the way everyone builds processors.

## ■ E

– Repeated a small datapath


– More "dark silicon" than active

## ■ Observations in the wild

– Flat frequency curve

– "Turbo Mode"

– Increasing cache/processor ratio






# Utilization Wall: Dark Implications for Multicore

Spectrum of tradeoffs  
between # of cores and  
frequency



Example:

65 nm  $\rightarrow$  32 nm ( $S = 2$ )

4 cores @ 1.8 GHz




65 nm



2x4 cores @ 1.8 GHz  
(8 cores dark, 8 dim)

*(Industry's Choice)*



4 cores @ 2x1.8 GHz  
(12 cores dark)

32 nm


# What do we do with dark silicon?

- Goal: Leverage dark silicon to scale the utilization wall
- Insights:
  - Power is now more expensive than area
  - Specialized logic can improve energy efficiency (10–1000x)
- Our approach:
  - Fill dark silicon with specialized cores to save energy on common applications
  - Provide focused reconfigurability to handle evolving workloads

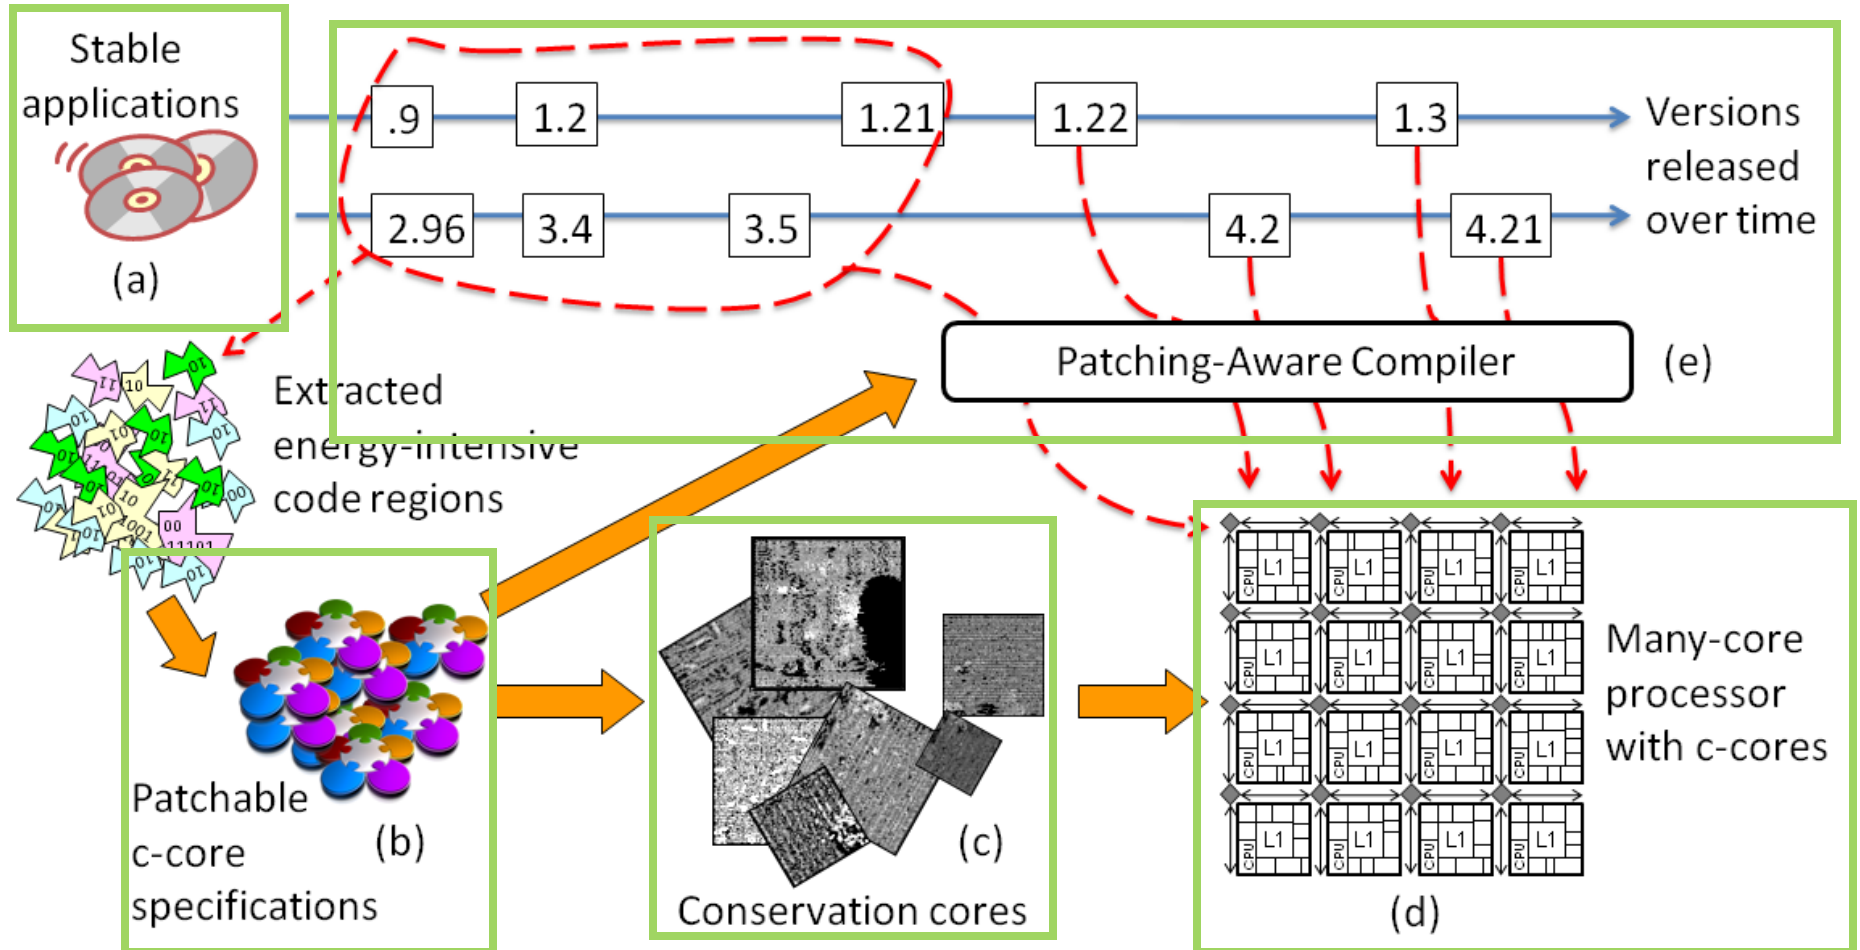
# Conservation Cores

*"Conservation Cores: Reducing the Energy of Mature Computations," Venkatesh et al., ASPLOS '10*

- Specialized circuits for reducing energy
  - Automatically generated from hot regions of program source code
  - Patching support future-proofs the hardware
- Fully-automated toolchain
  - Drop-in replacements for code
  - Hot code implemented by c-cores, cold code runs on host CPU
  - HW generation/SW integration
- Energy-efficient
  - Up to 18x for targeted hot code



# The C-core Life Cycle



# Outline


- Utilization wall and dark silicon
- GreenDroid
- Conservation cores
- GreenDroid energy savings
- Conclusions

# Emerging Trends

The *utilization wall* is exponentially worsening the dark silicon problem.


Specialized architectures are receiving more and more attention because of energy efficiency.

*Mobile application processors* are becoming a dominant computing platform for end users.



# Mobile Application Processors Face the Utilization Wall


- The evolution of mobile application processors mirrors that of microprocessors
- Application processors face the utilization wall
  - Growing performance demands
  - Extreme power constraints



# Android™



- Google's OS + app. environment for mobile devices
- Java applications run on the Dalvik virtual machine
- Apps share a set of libraries (libc, OpenGL, SQLite, etc.)

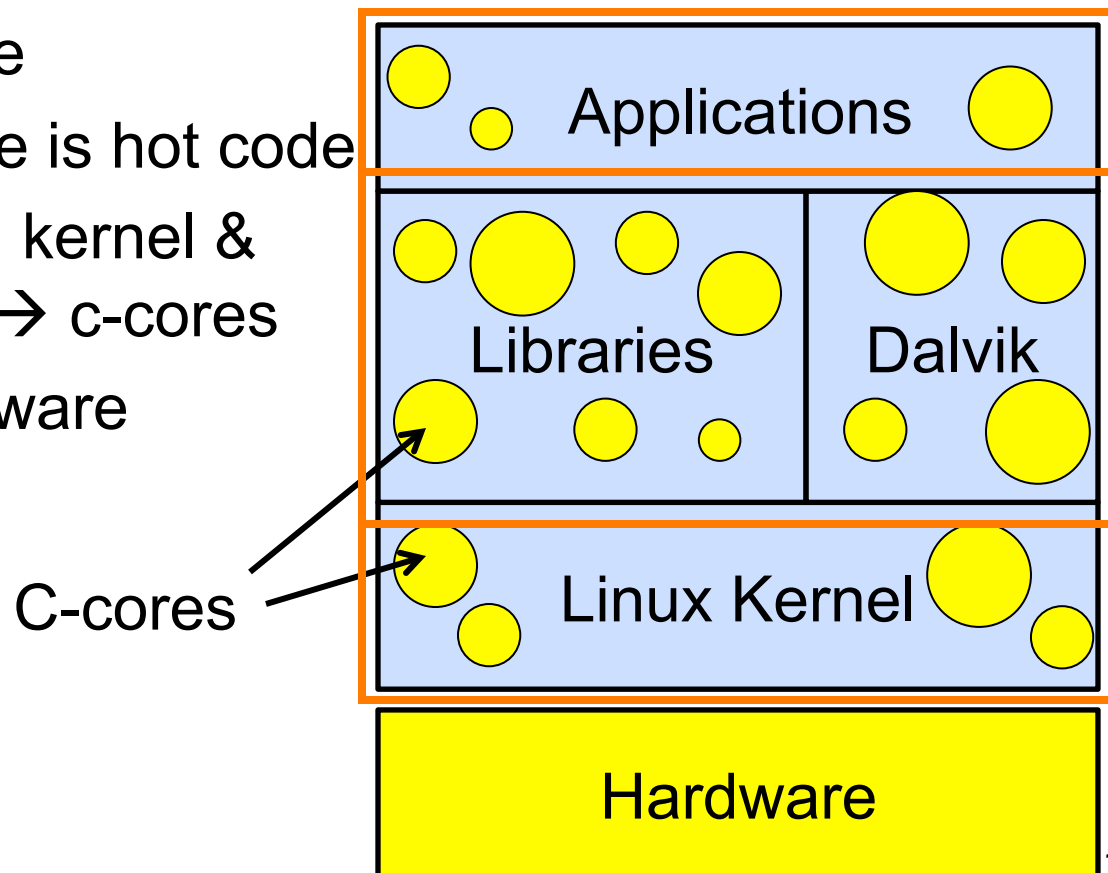




# Applying C-cores to Android




- Android is well-suited for c-cores
  - Core set of commonly used applications
  - Libraries are hot code
  - Dalvik virtual machine is hot code
  - Libraries, Dalvik, and kernel & application hotspots → c-cores
  - Relatively short hardware replacement cycle




# Android Workload Profile

- Profiled common Android apps to find the hot spots, including:
  - Google: Browser, Gallery, Mail, Maps, Music, Video
  - Pandora
  - Photoshop Mobile
  - Robo Defense game
- Broad-based c-cores
  - 72% code sharing
- Targeted c-cores
  - 95% coverage with just 43,000 static instructions (approx. 7 mm<sup>2</sup>)




# GreenDroid: Applying Massive Specialization to Mobile Application Processors




# GreenDroid Tiled Architecture

- Tiled lattice of 16 cores
- Each tile contains
  - 6-10 Android c-cores (~125 total)
  - 32 KB D-cache (shared with CPU)
  - MIPS processor
    - 32 bit, in-order, 7-stage pipeline
    - 16 KB I-cache
    - Single-precision FPU
  - On-chip network router




# GreenDroid Tile Floorplan

- 1.0 mm<sup>2</sup> per tile
- 50% C-cores
- 25% D-cache
- 25% MIPS core, I-cache, and on-chip network



# GreenDroid Tile Skeleton

- 45 nm process
- 1.5 GHz
- ~30k instances
- Blank space is filled with a collection of c-cores
- Each tile contains different c-cores



# Outline

- Utilization wall and dark silicon
- GreenDroid
- Conservation cores
- GreenDroid energy savings
- Conclusions

# Constructing a C-core

- C-cores start with source code
  - Can be irregular, integer programs
  - Parallelism-agnostic
- Supports almost all of C:
  - Complex control flow  
e.g., goto, switch, function calls
  - Arbitrary memory structures  
e.g., pointers, structs, stack, heap
  - Arbitrary operators  
e.g., floating point, divide
  - Memory coherent with host CPU

```
sumArray(int *a, int n)
{
    int i = 0;
    int sum = 0;

    for (i = 0; i < n; i++) {
        sum += a[i];
    }

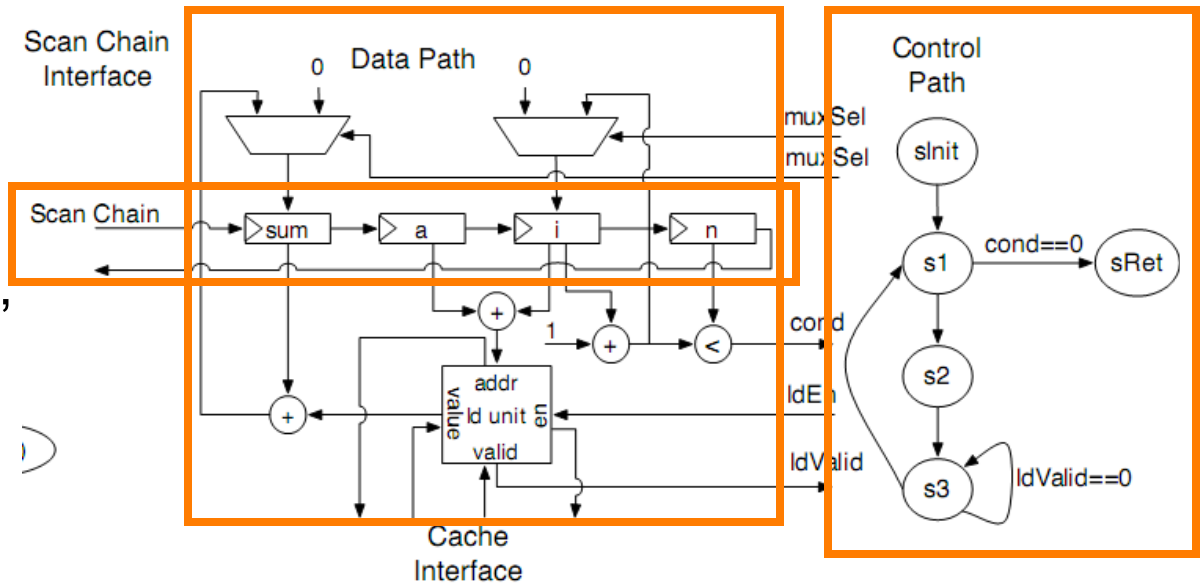
    return sum;
}
```



# Constructing a C-core


## ■ Compilation

- C-core selection
- SSA, infinite register, 3-address code
- Direct mapping from CFG and DFG
- Scan chain insertion



## ■ Verilog → Place & Route

- 45 nm process
- Synopsys CAD flow
  - Synthesis
  - Placement
  - Clock tree generation
  - Routing



# C-cores Experimental Data

- We automatically built 21 c-cores for 9 "hard" applications


- 45 nm TSMC

- Vary in size from 0.10 to 0.25 mm<sup>2</sup>

- Frequencies from 1.0 to 1.4 GHz


Application	# C-cores	Area (mm <sup>2</sup> )	Frequency (MHz)
bzip2	1	0.18	1235
cjpeg	3	0.18	1451
djpeg	3	0.21	1460
mcf	3	0.17	1407
radix	1	0.10	1364
sat solver	2	0.20	1275
twolf	6	0.25	1426
viterbi	1	0.12	1264
vpr	1	0.24	1074

# C-core Energy Efficiency: Non-cache Operations




- Up to 18x more energy-efficient (13.7x on average), compared to running on the MIPS processor

# Where do the energy savings come from?



MIPS baseline  
91 pJ/instr.




C-cores  
8 pJ/instr.

# Supporting Software Changes

- Software may change – HW must remain usable
  - C-cores unaffected by changes to cold regions
- Can support any changes, through *patching*
  - Arbitrary insertion of code – software exception mechanism
  - Changes to program constants – configurable registers
  - Changes to operators – configurable functional units
- Software exception mechanism
  - Scan in values from c-core
  - Execute in processor
  - Scan out values back to c-core to resume execution

# Patchability Payoff: Longevity

- Graceful degradation
  - Lower initial efficiency
  - Much longer useful lifetime
- Increased viability
  - With patching, utility lasts ~10 years for 4 out of 5 applications
  - Decreases risks of specialization




# Outline

- Utilization wall and dark silicon
- GreenDroid
- Conservation cores
- GreenDroid energy savings
- Conclusions

# GreenDroid: Energy per Instruction

- More area dedicated to c-cores yields higher execution coverage and lower energy per instruction (EPI)



- 7 mm² of c-cores provides:
  - 95% execution coverage
  - 8x energy savings over MIPS core



# What kinds of hotspots turn into GreenDroid c-cores?

C-core	Library	# Apps	Coverage (est., %)	Area (est., mm <sup>2</sup> )	Broad-based
→ dvmInterpretStd	libdvm	8	10.8	0.414	Y
→ scanObject	libdvm	8	3.6	0.061	Y
→ S32A_D565_Opaque_Dither	libskia	8	2.8	0.014	Y
→ src_aligned	libc	8	2.3	0.005	Y
→ S32_opaque_D32_filter_DXDY	libskia	1	2.2	0.013	N
→ less_than_32_left	libc	7	1.7	0.013	Y
→ cached_aligned32	libc	9	1.5	0.004	Y
.plt	<many>	8	1.4	0.043	Y
→ memcpy	libc	8	1.2	0.003	Y
→ S32A_Opaque_BlitRow32	libskia	7	1.2	0.005	Y
→ ClampX_ClampY_filter_affine	libskia	4	1.1	0.015	Y
DiagonalInterpMC	libomx	1	1.1	0.054	N
→ blitRect	libskia	1	1.1	0.008	N
calc_sbr_synfilterbank_LC	libomx	1	1.1	0.034	N
inflate	libz	4	0.9	0.055	Y
...	...	...	...	...	...

# GreenDroid: Projected Energy

Aggressive mobile application processor  
(45 nm, 1.5 GHz) 91 pJ/instr.

GreenDroid c-cores 8 pJ/instr.

GreenDroid c-cores + cold code (est.) 12 pJ/instr.

- GreenDroid c-cores use 11x less energy per instruction than an aggressive mobile application processor
- Including cold code, GreenDroid will still save ~7.5x energy

# Project Status

## ■ Completed

- Automatic generation of c-cores from source code to place & route
- Cycle- and energy-accurate simulation (post place & route)
- Tiled lattice, placed and routed
- FPGA emulation of c-cores and tiled lattice

## ■ Ongoing work

- Finish full system Android emulation for more accurate workload modeling
- Finalize c-core selection based on full system Android workload model
- Timing closure and tapeout

# GreenDroid Conclusions

- The utilization wall forces us to change how we build hardware
- Conservation cores use dark silicon to attack the utilization wall
- GreenDroid will demonstrate the benefits of c-cores for mobile application processors
- We are developing a 45 nm tiled prototype at UCSD

# GreenDroid: A Mobile Application Processor for a Future of Dark Silicon

---

Nathan Goulding, Jack Sampson, Ganesh Venkatesh,  
Saturnino Garcia, Joe Auricchio, Jonathan Babb<sup>+</sup>,  
Michael B. Taylor, and Steven Swanson

*Department of Computer Science and Engineering,  
University of California, San Diego*

<sup>+</sup> CSAIL, Massachusetts Institute of Technology

# Backup Slides

# Automated Measurement Methodology

- C-core toolchain
  - Specification generator
  - Verilog generator
- Synopsys CAD flow
  - Design Compiler
  - IC Compiler
  - 45 nm library
- Simulation
  - Validated cycle-accurate c-core modules
  - Post-route gate-level simulation
- Power measurement
  - VCS + PrimeTime

